



≡ Overview > Introduction

## Overview

# Introduction

📄 Copy page



Chroma is the open-source data infrastructure for AI. It comes with everything you need to get started built-in.

Chroma gives you everything you need for retrieval: store embeddings with metadata, search with dense and sparse vectors, filter by metadata, and retrieve across text, images, and more.

## What Chroma Offers

### Document Storage

Store documents and metadata.

### Embeddings

Use any embedding model. OpenAI, Cohere, Hugging Face, sentence-transformers, and more.

### Vector Search

Dense, sparse, and hybrid search. Query by similarity and combine multiple search strategies.

### Full-Text & Regex Search

### Metadata Filtering

Filter results at query time by metadata conditions.

### Multi-Modal Retrieval

Index and search images, audio, and



&gt;

## Quickstart

### Getting Started with the Chroma SDK

Create a self-hosted or cloud database and add data to it using the Chroma SDK.

### Create a Chroma Cloud Database

Create a scalable, zero-ops Chroma Cloud database to store your AI data.



## Example Projects

### Agentic Search

Build agents that iteratively search and refine results for complex queries.

### Code Search

Index codebases to power coding agents using AST-aware chunking.



## Open Source

Chroma is licensed under [Apache 2.0](#). Run it locally, self-host, or use [Chroma Cloud](#) for a managed, serverless experience.

Was this page helpful?

 Yes

 No

 Suggest edits



[Enterprise](#)

[Pricing](#)

[Changelog](#)